

MapReduce and Spark

MIE451
Scott Sanner

Slides are taken from Liping Liu and Ch. 2 of the book Mining of Massive Datasets: <http://www.mmms.org/>

The big data challenge

- Google Example
 - 20+ billion web pages x 20KB = 400+ TB
 - 1 computer reads 30-35 MB/sec from disk
 - ~4 months to read the web
 - ~1,000 hard drives to store the web
 - Takes even more to do something useful with the data!
- With “small data”, we still want to make the program faster with parallel computing


2

Distributed computation: overview

- **Today, a standard architecture for such problems is emerging:**
 - Cluster of commodity Linux nodes
 - Commodity network (ethernet) to connect them
- **Challenges:**
 - How to distribute computation?
 - Distributed/parallel programming is hard
 - Machines fail:
 - One server may stay up 3 years (1,000 days)
 - If you have 1,000 servers, expect to loose 1/day
- **Map-reduce** addresses all of the above

3

Cluster Architecture



Each rack contains 16-64 nodes

In 2011 it was guestimated that Google had 1M machines, <http://bit.ly/Shh0RO>

Distribute computation

- **Issue: Copying data over a network takes time**
- **Idea:**
 - Bring computation close to the data
 - Store files multiple times for reliability
- **Map-reduce addresses these problems**
 - Elegant way to work with big data
 - **Programming model**
 - Map-Reduce
 - **Storage Infrastructure – File system**
 - Google: GFS. Hadoop: HDFS


5

MapReduce Programming

- Read data
- **Map:**
 - Extract something you care about
- **Sort and Shuffle:** Group by key
- **Reduce:**
 - Aggregate, summarize, filter or transform
- Write the result

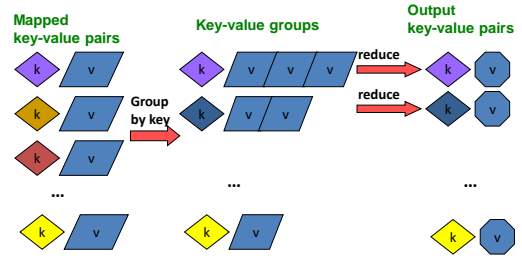
6

MapReduce: The Map Step



7

MapReduce: The Reduce Step



8

More Specifically

- **Input:** a set of key-value pairs
- Programmer specifies two methods:
 - **Map(line in data file)** $\rightarrow \langle k', v' \rangle^*$
 - Takes a line in data file and outputs **set** of key-value pairs
 - There is one Map call for every data file line
 - **Reduce($k', \langle v' \rangle^*$)** $\rightarrow \langle k', v'' \rangle$
 - All values v' with same key k' are reduced together and processed in v'' order
 - There is one Reduce function call per unique key k'

9

Word count example

- We have a huge text document
- Count the number of times each distinct word appears in the file
 - File too large for memory, but all $\langle \text{word}, \text{count} \rangle$ pairs fit in memory
- **Sample application:**
 - Analyze web server logs to find popular URLs

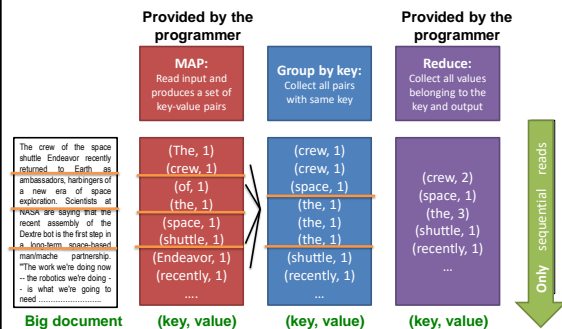
10

Task: Word Count

- A method that works on a **small document**
 - `words(doc.txt) | sort | uniq -c`
 - where `words` takes a file and outputs the words in it, one per a line
- This method captures the essence of **MapReduce**
 - Great thing is that it is naturally **parallelizable**

11

MapReduce: Word Counting



12

Word Count Using MapReduce

```
map(key, value):
// key: document name; value: text of the document
for each word w in value:
    emit(w, 1)

reduce(key, values):
// key: a word; value: an iterator over counts
result = 0
for each count v in values:
    result += v
emit(key, result)
```

13

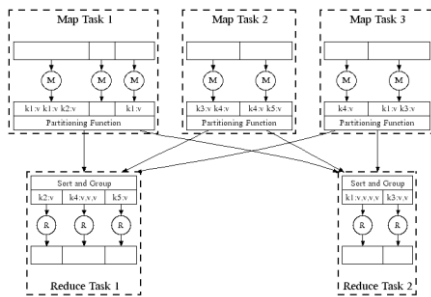
Map-Reduce: Environment

Map-Reduce environment takes care of:

- Partitioning the input data
- Scheduling the program's execution across a set of machines
- Performing the **group by key** step
- Handling machine **failures**
- Managing required inter-machine **communication**

14

Map-Reduce: In Parallel




All phases are distributed with many tasks doing the work

15

Map-Reduce

- **Programmer specifies:**
 - Map and Reduce and input files
- **Workflow:**
 - Read file lines as key-value-pairs
 - **Map** transforms input kv-pairs into a new set of k'v'-pairs
 - Sorts & Shuffles the k'v'-pairs
 - All k'v'-pairs with a given k' are sent to the same **reduce**
 - **Reduce** processes all k'v'-pairs grouped by key into new k''v''-pairs
 - Write the resulting pairs to files
- All phases are distributed with many tasks doing the work



16


Data Flow

- **Input and final output are stored on a distributed file system (FS):**
 - Scheduler tries to schedule map tasks “close” to physical storage location of input data
- **Intermediate results are stored on local FS of Map and Reduce workers**
- **Output is often input to another MapReduce task**

17

Refinement: Combiners

- Often a Map task will produce many pairs of the form $(k, v_1), (k, v_2), \dots$ for the same key k
 - E.g., popular words in the word count example
- **Can save network time by pre-aggregating values in the mapper:**
 - $\text{combine}(k, \text{list}(v_i)) \rightarrow v_2$
 - Combiner is usually same as the reduce function
- Works only if reduce function is commutative and associative,
 - Consider sum, max, mean and stddev




18

Refinement: Combiners

- **Back to our word counting example:**

- Combiner combines the values of all keys of a single mapper (single machine):



- Much less data needs to be copied and shuffled!

19

Refinements: Backup Tasks

- **Problem**

- Slow workers significantly lengthen the job completion time:
 - Other jobs on the machine
 - Bad disks
 - Weird things

- **Solution**

- Near end of phase, spawn backup copies of tasks
 - Whichever one finishes first “wins”

- **Effect**

- Dramatically shortens job completion time

20

Coordination: Master

- **Master node takes care of coordination:**

- **Task status:** (idle, in-progress, completed)
- **Waiting tasks** get scheduled as workers become available
- When a map task completes, it sends the master the location and sizes of its R intermediate files, one for each reducer
- Master pushes this info to reducers

- Master pings workers periodically to detect failures

21

Storage Infrastructure

- **Problem:**

- If nodes fail, how to store data persistently?

- **Answer:**

- **Distributed File System:**

- Provides global file namespace

- **Typical usage pattern**

- Huge files (100s of GB to TB)
- Data is rarely updated in place
- Reads and appends are common

22

Distributed File System

- **Chunk servers**

- File is split into contiguous chunks
- Typically each chunk is 16-64MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

- **Master node**

- a.k.a. Name Node in Hadoop's HDFS
- Stores metadata about where files are stored
- Might be replicated

- **Client library for file access**

- Talks to master to find chunk servers
- Connects directly to chunk servers to access data

23

Distributed File System

- **Reliable distributed file system**

- Data kept in “chunks” spread across machines

- Each chunk **replicated** on different machines

- Seamless recovery from disk or machine failure



Bring computation directly to the data!

Chunk servers also serve as compute servers

24