

Machine Learning: Big Picture

Scott Sanner
U. of Toronto, MIE

Machine Learning

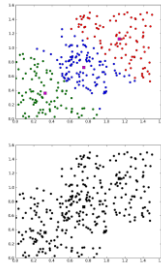
- You have a data set $D = \{(x_i, y_i)\}$
 - x_i is a feature vector, y_i are labels
 - Data may be partially specified
- You want to learn $y = f(x)$ from D
 - More precisely, you want to minimize some error $E(f, w, D)$

Function class Function parameters

- Majority of problems are either
 - Classification:** y is discrete
 - Regression:** y is continuous

Supervised vs. Unsupervised

- $D = \{(x_i, y_i)\}$
- Supervised:** y_i observed
 - Classification for discrete y_i
 - Regression for continuous y_i
- Unsupervised:** y_i not observed
 - Learning method has to assign y_i
 - E.g., clustering via K-means



Inductive Bias

- Let's avoid making assumptions about f**
 - Assume for simplicity that $D = \{(x_i, y_i)\}$ is noise free
 - x_i 's in D only cover small subset of input space x
- What's the best we can do?**
 - If we've seen $x=x_i$ report $y=y_i$
 - If we have not seen $x=x_i$, can't say anything (no assumptions)
- This is called rote learning... boring, eh?**
 - Key idea: *you can't generalize to unseen data w/o assumptions!*
- Thus, key to ML is generalization**
 - To generalize, ML algorithm *must* have some **inductive bias**
 - Bias usually in the form of a **restricted hypothesis space**
 - Important to understand restrictions (and whether appropriate)

Parametric vs. Non-parametric

- Parametric:** has parameters
 - Most Probabilistic Approaches
 - Gaussians
 - Bernoulli / Binomial / Multinomial
 - Graphical Models
 - Linear Regression / Classifiers
 - SVM with linear kernel
- Non/semi-parametric:** data oriented
 - Neighbor-based approaches
 - (K-)nearest Neighbor
 - Parzen Windows
 - SVM with RBF kernel

What assumptions?

Data is generated by given distribution.
Linearly separable, error assumptions (e.g., Gaussian), etc.

What assumptions?

Encoded in distance function & K / width.
Smoothness... no abrupt changes!

Linear vs. Non-linear

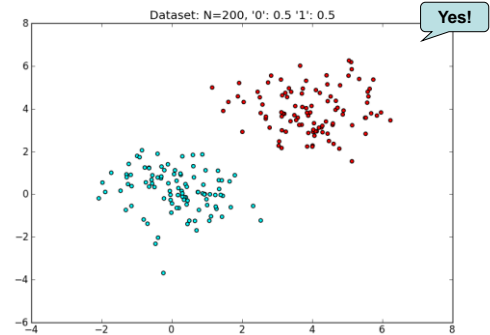
- $y = f(x, w)$
 - x is your input vector
 - w is your parameter vector (weights)
- Which f is linear in w ?**
i.e., $f(x, w) = \langle w, x \rangle$ (assume $x_0 = 1$)
 - $f_1(x) = w_1 x_1 + w_2 x_2$ ✓
 - $f_2(x) = w_1 x_1^2 + w_2 x_1 x_2 + w_3 x_2^2$ ✓
 - $f_3(x) = w_1 x_1 + w_2 w_3 x_2 + w_3^2 x_3$ ✗

Any transformation of input x maintains linear function in w !

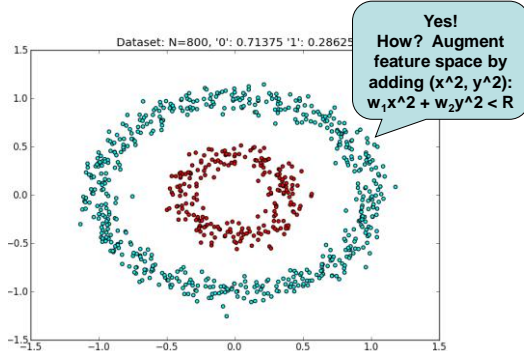
Your Linear Fun. Approx. Toolbox

- **Classification**
 - Naïve Bayes (simple)
 - Logistic Regression (better than NB for dependent features)
 - Perceptron (didactic, rarely used in practice)
 - SVM and Kernel Methods (very powerful)
- **Regression**
 - Linear Regression (closed-form solution)
 - SVR and Kernel Methods (very powerful)
- **Key Advantage**
 - All of above lead to convex optimization problems
→ global optima will be found.

Classes Linearly Separable?



Classes Linearly Separable?



How Powerful are Linear Models?

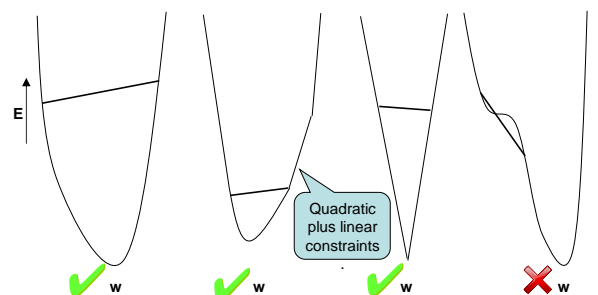
- **Short answer:**
 - $y = \langle w, x \rangle = \sum_i w_i x_i$ is surprisingly powerful
 - Especially if x transformed: $x \rightarrow \Phi(x)$
 - $y = \langle w, \Phi(x) \rangle$
- **Can use expressive feature spaces $\Phi(x)$**
 - If data \ll features, beware of **overfitting**
 - More on overfitting later
 - Allows to fit virtually any nonlinear function
 - Optimization problem is still convex – **solve optimally!**

This is Worth Investigating...

- **Seems counterintuitive...**
 - Fit crazy non-linear functions & find global optimum?
 - WTF? **Why's that fine?**
- **Think about abstracted problem...**
 - $E(w) = E(f, w, D)$ (b/c f, D fixed)
 - We change the weights w , we get different $E(w)$.
 - Which setting of weights optimizes $E(w)$?
- **Question: how does $E(w)$ look w.r.t. weights?**
 - Linear regression: linear f & SSE, $E(w)$ looks **quadratic**
 - In general, can show Hessian (2nd derivative matrix) is positive semidefinite \Rightarrow **convex**...

Convexity

- Graphically... which of these is convex?



Empirical Risk Minimization (ERM)

- A **general framework** for function approximation
- Minimize $E'(\mathbf{w}) = \text{Loss}(\mathbf{w}) + C \cdot \text{Regularizer}(\mathbf{w})$
 - Critical to avoid overfitting
- Loss functions penalize errors in different ways, e.g.,
 - Sum of squared error (SSE) → **Linear Regression**
 - Hinge loss → Why useful?
 - ϵ -insensitive loss → Why useful?
- Regularizer expresses preference on \mathbf{w} , e.g.,
 - $\|\mathbf{w}\|_2$: assumes Gaussian prior (prefers small weights)
 - $\|\mathbf{w}\|_1$: can encourage sparsity
 - $\mathbf{w} \cdot \log \mathbf{w}$: maximizes entropy for prob. interpretation of \mathbf{w}
- Many $E'(\mathbf{w})$ possibilities are convex for linear f !

The Joy of Convex(ity)

All of these losses are convex for linear f :

Table 1: Scalar loss functions and their derivatives, depending on $f := \langle \mathbf{w}, \mathbf{x} \rangle$, and y .

	Loss $\ell(f, y)$	Derivative $\ell'(f, y)$
Hinge [20]	$\max(0, -yf)$	0 if $yf \geq 0$ and $-y$ otherwise
Squared Hinge [26]	$\frac{1}{2} \max(0, -yf)^2$	0 if $yf \geq 0$ and f otherwise
Soft Margin [4]	$\max(0, 1 - yf)$	0 if $yf \geq 1$ and $-y$ otherwise
Squared Soft Margin [10]	$\frac{1}{2} \max(0, 1 - yf)^2$	0 if $yf \geq 1$ and $f - y$ otherwise
Exponential [14]	$\exp(-yf)$	$-y \exp(-yf)$
Logistic [13]	$\log(1 + \exp(-yf))$	$-y / (1 + \exp(yf))$
Novelty [32]	$\max(0, 1 - f)$	0 if $f \geq 0$ and -1 otherwise
Least mean squares [43]	$\frac{1}{2}(f - y)^2$ ← linear regression	$f - y$
Least absolute deviation	$ f - y $	$\text{sgn}(f - y)$
Quantile regression [27]	$\max(\tau(f - y), (1 - \tau)(y - f))$	τ if $f > y$ and $\tau - 1$ otherwise
ϵ -insensitive [41]	$\max(0, f - y - \epsilon)$	0 if $ f - y \leq \epsilon$ and $\text{sgn}(f - y)$ otherwise
Huber's robust loss [31]	$\frac{1}{2}(f - y)^2$ if $ f - y < 1$, else $ f - y - \frac{1}{2}$	$f - y$ if $ f - y \leq 1$, else $\text{sgn}(f - y)$
Poisson regression [16]	$\exp(f) - yf$	$\exp(f) - y$

Table 2: Vectorial loss functions and their derivatives, depending on the vector $f := Wx$ and on y .

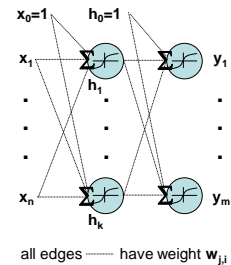
	Loss	Derivative
Soft Margin [38]	$\max_{y'} (f_{y'} - f_y + \Delta(y, y'))$	$e_{y^*} - e_y$, where y^* is the argmax of the loss
Scaled Soft Margin [40]	$\max_{y'} \Delta^s(y, y') (f_{y'} - f_y + \Delta(y, y'))$	$\Delta^s(y, y') (e_{y^*} - e_y)$, where y^* is the argmax of the loss
Softmax [14]	$\log \sum_{y'} \exp(f_{y'}) - f_y$	$\sum_{y'} e_{y'} \exp(f_{y'}) / \sum_{y'} \exp(f_{y'}) - e_y$
Multivariate Regression	$\frac{1}{2}(f - y)^T M(f - y)$ where $M \succeq 0$	$M(f - y)$

Non-convexity

- Convexity was the rage from 2000-2010
- Non-convex approaches like neural networks could not guarantee optimal learning
 - Therefore no one used them
 - No one bothered to compare to them
 - Reviewers rejected papers on the topic
 - Until some folks recently tried using them in combination with Big Data + Many Layers + GPUs
- And they led to massive improvements!

Artificial Neural Networks

- Neural Net:** non-linear weighted combination of *shared* sub-functions
- Backpropagation:** to minimize SSE, train weights using *gradient descent* and *chain rule*
- Deep nets:** have more than one hidden layer
<http://playground.tensorflow.org/>

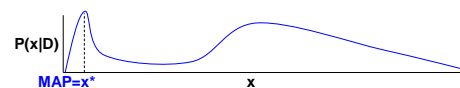


Being Bayesian

- Many (philosophical) interpretations of what it means to be Bayesian
 - All share a common characteristic
- Bayesian = maintaining a distribution over the most likely values of a (random) variable**
 - Variable can be any quantity of interest
 - A location (\mathbf{x}, \mathbf{y}) for tracking
 - Parameters \mathbf{w} of a linear classifier

Why be Bayesian? Risk!

- Robot has Bayesian belief $P(\mathbf{x}|\mathbf{D})$ over position \mathbf{x}**
 - \mathbf{D} consists of noisy range finder readings



- Associate Risk(x) with position \mathbf{x} (e.g., stairs!)**
 -
 - MAP Risk = $\text{Risk}(x^*) = 0$
 - Full Bayesian Risk = $\int_x \text{Risk}(x) p(x|\mathbf{D}) dx > 0$
- Which risk estimate would you use?

Overfitting

- In brief: fitting characteristics of training data that do not generalize to future test data
- Central problem in machine learning
 - Particularly problematic if $\#data \ll \#parameters$
 - ... don't have enough data to "identify" parameters

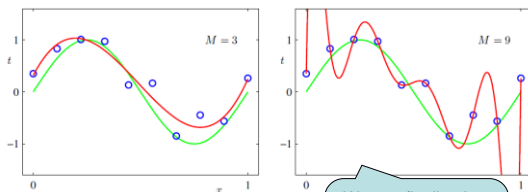
Overfitting in Classification

- Example: try to classify technology web pages {true, false}
 - Crawl Microsoft website as positive examples
 - Highly weighted features:
 - Bill Gates
 - support@microsoft.com
 - But do they generalize?

No, we've overfit to sampled-biased training data. Can we combat this?

Overfitting in Regression

- Green: Generate data (blue) plus noise
- Red: Fit a polynomial of degree M to the data



We can fit all points perfectly with a 9th degree polynomial... will its predictions generalize?

Combating Overfitting

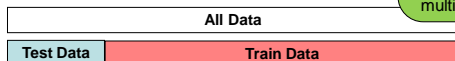
- Careful "unbiased" selection of training data helps
 - Test data should be from same distribution as train data
- But unbiased data is not easy, nor enough
 - There are always spurious correlations to overfit
 - Best way to fight overfitting is by **restricting the hypothesis space**
 - For linear classifiers, do feature selection
 - Tune "hyperparameters" to avoid overfitting
 - Naïve Bayes has smoothing hyperparameter
 - SVMs and Logistic Regression have "C"

How to fix problem on previous slide?

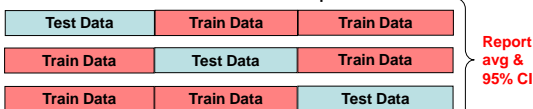
- Cross-validation is important for tuning**

(Cross) Validation

- Validation:**
 - Split your data into train and test



- K-fold Cross Validation:**
 - Reduce variance of test error estimate
 - 3-fold cross validation example



Sometimes a single random split may do unusually good or bad... better to average over multiple splits.

Repeated Random Sub-sampling Validation

- If have 1000 samples and need $k=100$
 - Then test case only has 10 samples
 - Too small!
- Instead, choose an $X\%$ train/test split k times:
 - For $X = 90$, randomly split data into 90% for train data and 10% test data
 - Take $k=100$ splits with test data containing 100 samples
 - Caveat: testing data overlaps, introduce correlations
 - Make sure to use a proper random permutation
 - `numpy.random.shuffle` or `numpy.random.permutation`

Nested Cross-Validation (NCV)

- Two levels of cross-validation

- 1st (Top) level CV

- For **average performance**
- Split data into {train, test}

Why NCV? Because must tune to fight overfitting, but cannot tune on the test set! This is cheating... why?

Can't tune on train data – need tuning to generalize to unseen data.

- 2nd (Bottom) level CV

- For **hyperparameter tuning**
- Split train into {train', validation}
- Choose hyperparameters that maximize avg performance on validation set, use in 1st level

So separate out validation from train in a CV sub-level to tune hyperparameters.

If takes too long, tune on single train/val split at 2nd level, but less robust.

Aside: CI Common Mistake

- When I ask for **avg \pm CIs**

- I want confidence interval (CIs) on the **avg**

- 67% CIs are roughly **avg $\pm \sigma/\sqrt{n}$**

- 95% CIs are roughly **avg $\pm 2\sigma/\sqrt{n}$**

How confident are you for $n=\infty$?

- σ/\sqrt{n} is standard error of the mean

– https://en.wikipedia.org/wiki/Standard_error#Standard_error_of_the_mean_2

- **95% CIs**: intervals for avg that should hold 95 out of 100 times if I rerun the randomized experiment

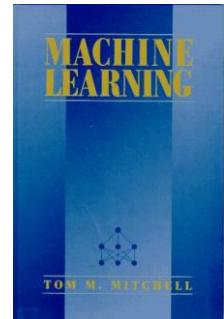
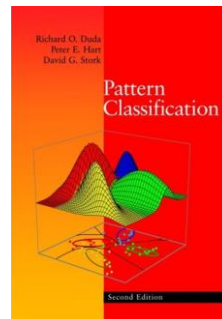
- Students often give me **avg $\pm \sigma$** (std deviation)

- Sample deviation, not confidence in sample mean

Model and Feature Selection

- Data set **D = {(x_i, y_i)}**
- Learn **y = f(x)** by minimizing **E(f, w, D)**
- Model Selection:
 - Which **f** to use
 - Linear / Non-linear
 - Parametric / Non- / Semi-parametric
 - Parameters within each model
- Feature Selection
 - Which **x** to use
 - Subset, transformed?
- How to choose between different models or feature sets
 - Can choose model / feature set with **lowest CV error**

Introductory Books to Consider



Chapters 1-6 in both books are useful.