


Sentiment Analysis

What is Sentiment Analysis?



Positive or negative movie review?

- unbelievably disappointing
- Full of zany characters and richly applied satire, and some great plot twists
- this is the greatest screwball comedy ever filmed
- It was pathetic. The worst part about it was the boxing scenes.



Google Product Search



HP Officejet 6500A Plus e-All-in-One Color Ink-jet - Fax / copier / printer / scanner
\$89 online, \$100 nearby ★★★★★ 377 reviews
September 2010 - Printer - HP - Inkjet - Office - Copier - Color - Scanner - Fax - 250 sh


Reviews

Summary - Based on 377 reviews


1 star 2 3 4 stars 5 stars

What people are saying

- ease of use "This was very easy to setup to four computers."
- value "Appreciate good quality at a fair price."
- setup "Overall pretty easy setup."
- customer service "I DO like honest tech support people."
- size "Pretty Paper weight."
- mode "Photos were fair on the high quality mode."
- colors "Full color prints came out with great quality."



Bing Shopping



HP Officejet 6500A E710N Multifunction Printer
Product summary Find best price Customer reviews Specifications Related items

\$121.53 - \$242.39 (14 stores)

Compare


Average rating ★★★★★ (144)

Most mentioned

- Performance (57)
- Ease of Use (54)
- Print Speed (43)
- Connectivity (39)
- More ▼ (31)

Show reviews by source

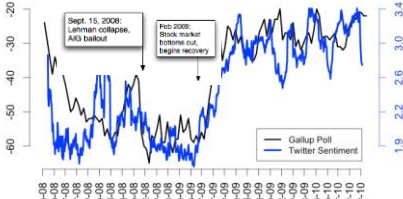
- Best Buy (140)
- CNET (6)
- Amazon.com (3)
- More ▼




Twitter sentiment versus Gallup Poll of Consumer Confidence

Brendan O'Connor, Ramnath Balasubramanyam, Bryan R. Routledge, and Noah A. Smith. 2010. From Tweets to Polls: Linking Text Sentiment to Public Opinion Time Series. In ICWSM-2010

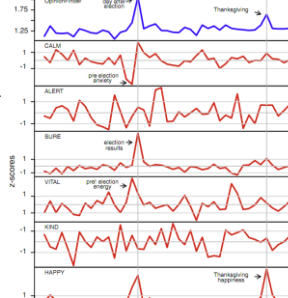
window = 15, $r = 0.804$

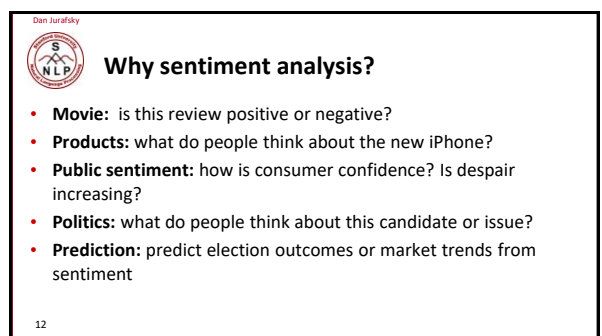

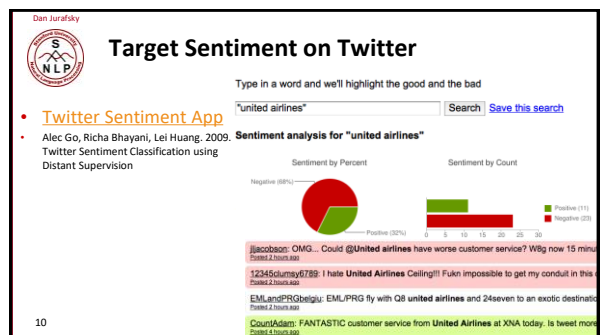








Twitter sentiment:

Johan Bollen, Huina Mao, Xiaojun Zeng. 2011. [Twitter mood predicts the stock market](#). Journal of Computational Science 2:1, 1-8. 10.1016/j.jocs.2010.12.007.








Dien Krafky

Scherer Typology of Affective States

- **Motion:** related organically synchronized ... evaluation of a major event
 - *angry, sad, joyful, fearful, ashamed, proud, elated*
- **Mood:** diffuse non-caused low-intensity long-duration change in subjective feeling
 - *cheerful, gloomy, irritable, listless, depressed, buoyant*
- **Interpersonal stances:** affective stance toward another person in a specific interaction
 - *friendly, flirtatious, distant, cold, warm, supportive, contemptuous*
- **Attitudes:** enduring, affectively colored beliefs, dispositions towards objects or persons
 - *liking, loving, hating, valuing, desiring*
- **Personality traits:** stable personality dispositions and typical behavior tendencies
 - *nervous, anxious, reckless, morose, hostile, jealous*




Don Juralaby

Recall that public mood said to predict stock market

Scherer Typology of Affective States

- **Emotion:** formed organically synchronized ... evaluation of a major event
 - *angry, sad, joyful, fearful, ashamed, proud, elated*
- **Mood:** diffuse non-caused low-intensity long-duration change in subjective feeling
 - *cheerful, gloomy, irritable, listless, depressed, buoyant*
- **Interpersonal stances:** affective stance toward another person in a specific interaction
 - *friendly, flirtatious, distant, cold, warm, supportive, contemptuous*
- **Attitudes:** enduring, affectively colored beliefs, dispositions towards objects or persons
 - *liking, loving, hating, valuing, desiring* Majority of review sentiment focuses on this
- **Personality traits:** stable personality dispositions and typical behavior tendencies
 - *nervous, anxious, reckless, morose, hostile, jealous*

Dan Jurafsky




Sentiment Analysis

- Sentiment analysis is the detection of **attitudes**
“enduring, affectively colored beliefs, dispositions towards objects or persons”
 1. **Holder (source)** of attitude
 2. **Target (aspect)** of attitude
 3. **Type** of attitude
 - From a set of types
 - *Like, love, hate, value, desire, etc.*
 - Or (more commonly) simple weighted **polarity**:
 - *positive, negative, neutral*, together with *strength*
 4. **Text** containing the attitude
 - Sentence or entire document

15


Don Jurafsky



Sentiment Analysis

- Simplest task:
 - Is the attitude of this text positive or negative?
- More complex:
 - Rank the attitude of this text from 1 to 5
- Advanced:
 - Detect the target, source, or complex attitude types

Don Jurafsky




Sentiment Analysis

- Simplest task:
 - Is the attitude of this text positive or negative?
- More complex:
 - Rank the attitude of this text from 1 to 5
- Advanced:
 - Detect the target, source, or complex attitude types

17

[illegible]




Stanford
NLP

Sentiment Classification in Movie Reviews

Bo Pang, Lillian Lee, and Shivakumar Vaithyanathan. 2002. Thumbs up? Sentiment Classification using Machine Learning Techniques. EMNLP-2002, 79–86.


Bo Pang and Lillian Lee. 2004. A Sentimental Education: Sentiment Analysis Using Subjectivity Summarization Based on Minimum Cuts. ACL, 271-278

- Polarity detection:
 - Is an IMDB movie review positive or negative?
- Data: *Polarity Data 2.0*:
 - <http://www.cs.cornell.edu/people/pabo/movie-review-data>



Stanford University
Natural Language Processing

IMDB data in the Pang and Lee database




when `_star` was `_came` out some twenty years ago, the image of traveling throughout the stars has become a commonplace image . [...]

when han solo goes light speed , the stars change to bright lines , going towards the viewer in lines that converge at an invisible point .

cool .

...
 _ocean sky _ offers a much simpler image—that of a single white dot , traveling horizontally across the night sky . [...]




“ snake eyes ” is the most aggravating kind of movie : the kind that shows so much potential then becomes unbelievably disappointing .

it's not just because this is a brian depalma film , and since he's a great director and one who's films are always greeted with at least some fanfare .


and it's not even because this was a film starring nicolas cage and since he gives a bravura performance , this film is hardly worth his talents .

© Dan Jurafsky



Baseline Algorithm (adapted from Pang and Lee)

- Tokenization
- Feature Extraction
- Classification using different classifiers
 - Naïve Bayes
 - MaxEnt
 - SVM



Stanford
Natural Language Processing

Sentiment Tokenization Issues

- Deal with HTML and XML markup
- Twitter mark-up (names, hash tags)
- Capitalization (preserve for

words in all caps)


```

[<>]?          # optional hat/brow
[:;#]?        # eyes
[\~o~\^*']?   # optional nose
[\]\]\ ([dDpP/\:;:]\{8\}\|)
              # mouth

[\]\]\ ([dDpP/\:;:]\{8\}\|)
              ### reverse orientation
[\~o~\^*']?   # mouth
[\~o~\^*']?   # eyes
[:;#]?        # optional hat/brow
[<>]?


```
- Phone numbers, dates
- Emoticons
- Useful code:
 - [Christopher Potts sentiment tokenizer](#)
 - [Brendan O'Connor twitter tokenizer](#)

Dan Jurafsky

 **Extracting Features for Sentiment Classification**

- How to handle negation
 - I **didn't** like this movie
 - vs
 - I really like this movie
- Which words to use?
 - Only adjectives
 - All words
 - All words turns out to work better, at least on this data

23



Dan Jurafsky

Negation

Dai, Sanjiv and Mike Chen. 2001. Yahoo!: Extracting market sentiment from stock message boards. In Proceedings of the Asia Pacific Finance Association Annual Conference (APFA).

Bo Pang, Lillian Lee, and Shivakumar Vaithyanathan. 2002. Thumbs up? Sentiment Classification using Machine Learning Techniques: EMNLP-2002, 79–86.

Add NOT_ to every word between negation and following punctuation:

didn't like this movie , but I

↓

didn't NOT_like NOT_this NOT_movie but I

24



Binarized (Boolean) features

- Intuition:
 - For sentiment (and probably for other text classification domains)
 - Word occurrence may matter more than word frequency
 - The occurrence of the word *fantastic* tells us a lot
 - The fact that it occurs 5 times may not tell us much more.

25



Classification Results

- SVM Classifier achieves 92.1% accuracy on positive/negative!
- But need to train a classifier per domain
 - Scary movie = good!
 - Scary hotel = bad!
 - Hotel with "thin walls"?

26



Classifiers don't capture everything

- Subtlety:
 - Perfume review in *Perfumes: the Guide*:
 - "If you are reading this because it is your darling fragrance, please wear it at home exclusively, and tape the windows shut."
 - Dorothy Parker on Katherine Hepburn
 - "She runs the gamut of emotions from A to B"

27



Classifiers don't capture everything

- Word interactions:
 - Raising taxes = bad
 - Raising salaries = good
 - Lowering taxes = good
 - Lowering salaries = bad
- A solution? Adjective*Noun where...
 - Raising=+1, Lowering=-1, Taxes=-1, Salaries=+1
 - Also supports Adverb*Adjective: very(+2) happy (+1), very (+2) sad (-1)

28



Thwarted Expectations and Ordering Effects

- "This film should be **brilliant**. It sounds like a **great** plot, the actors are **first grade**, and the supporting cast is **good** as well, and Stallone is attempting to deliver a good performance. However, it **can't hold up**."
- Well as usual Keanu Reeves is nothing special, but surprisingly, the **very talented** Laurence Fishbourne is **not so good** either, I was surprised.

29



Sarcasm

- Josef Stalin might enjoy this movie.
- As exciting as watching the grass grow.
- This movie should win flop of the year.
- I wondered whether I had checked into the Bates Motel.


30



Disagreements between polarity lexicons

Christopher Potts, [Sentiment Tutorial](#), 2011

	Opinion Lexicon	General Inquirer	SentiWordNet	LIWC
MPQA	33/5402 (0.6%)	49/2867 (2%)	1127/4214 (27%)	12/363 (3%)
Opinion Lexicon		32/2411 (1%)	1004/3994 (25%)	9/403 (2%)
General Inquirer			520/2306 (23%)	1/204 (0.5%)
SentiWordNet				174/694 (25%)
LIWC				



Analyzing the polarity of each word in IMDB

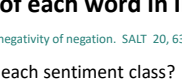
Potts, Christopher. 2011. On the negativity of negation. SALT 20, 636-659.

- How likely is each word to appear in each sentiment class?
- Count("bad") in 1-star, 2-star, 3-star, etc.
- But can't use raw counts:

- Instead, **likelihood**:


$$P(w|c) = \frac{f(w,c)}{\sum_w f(w,c)}$$

- Make them comparable between words
 - Scaled likelihood:** $\frac{P(w|c)}{P(w)}$




Counts of "bad" in IMDB

Actually what is this really?



Dan Jurafsky



Analyzing the polarity of each word in IMDB

Potts, Christopher. 2011. On the negativity of opinion. SALT 20, 636-659.

Category	Word	Count	Rating
POS	good	883,417	0.93
	amazing	103,509	0.17
	great	648,110	0.17
	awesome	47,142	0.27
NEG	good	26,447	0.09
	disappoint(ed)	118,498	0.13
	bad	368,273	0.04
	terrible	55,492	0.29

Open Jursafsky

Other sentiment feature: Logical negation

Potts, Christopher. 2011. On the negativity of negation. SALT 20, 636-659.

- Is logical negation (*no*, *not*) associated with negative sentiment?
- Potts experiment:
 - Count negation (*not*, *n't*, *no*, *never*) in online reviews
 - Regress against the review rating

40

**Potts 2011 Results:
More negation in negative sentiment**

IMDB (4,073,228 tokens) Five-star reviews (846,444 tokens)

Scaled likelihood $P(s|w)$

Sentiment Class	1	2	3	4	5	6	7	8	9	10
NEG	0.125	0.120	0.115	0.110	0.105	0.100	0.095	0.090	0.085	0.080
POS	0.085	0.085	0.085	0.085	0.085	0.085	0.085	0.085	0.085	0.085

Five-star reviews (846,444 tokens)

Sentiment Class	1	2	3	4	5
NEG	0.26	0.26	0.20	0.15	0.14
POS	0.19	0.19	0.19	0.19	0.19

44

45

46

47

48

Dan Jurafsky

Hatzivassiloglou & McKeown 1997

Step 4

- Clustering for partitioning the graph into two

49

Dan Jurafsky

Output polarity lexicon

- Positive
 - bold decisive disturbing generous good honest important large mature patient peaceful positive proud sound stimulating straightforward strange talented vigorous witty...
- Negative
 - ambiguous cautious cynical evasive harmful hypocritical inefficient insecure irrational irresponsible minor outspoken pleasant reckless risky selfish tedious unsupported vulnerable wasteful...

50

Dan Jurafsky

Output polarity lexicon

- Positive
 - bold decisive **disturbing** generous good honest important large mature patient peaceful positive proud sound stimulating straightforward **strange** talented vigorous witty...
- Negative
 - ambiguous **cautious** cynical evasive harmful hypocritical inefficient insecure irrational irresponsible minor **outspoken pleasant** reckless risky selfish tedious unsupported vulnerable wasteful...

51

Dan Jurafsky

Turney Algorithm

Turney (2002): Thumbs Up or Thumbs Down? Semantic Orientation Applied to Unsupervised Classification of Reviews

1. Extract a *phrasal lexicon* from reviews
2. Learn polarity of each phrase
3. Rate a review by the average polarity of its phrases

52

Dan Jurafsky

Extract two-word phrases with adjectives

First Word	Second Word	Third Word (not extracted)
JJ	NN or NNS	anything
RB, RBR, RBS	JJ	Not NN nor NNS
JJ	JJ	Not NN or NNS
NN or NNS	JJ	Nor NN nor NNS
RB, RBR, or RBS	VB, VBD, VBN, VBG	anything

53

Dan Jurafsky

How to measure polarity of a phrase?

- Positive phrases co-occur more with "excellent"
- Negative phrases co-occur more with "poor"
- But how to measure co-occurrence?

54



Pointwise Mutual Information

- **Mutual information** between 2 random variables X and Y

$$I(X, Y) = \sum_x \sum_y P(x, y) \log_2 \frac{P(x, y)}{P(x)P(y)}$$

- **Pointwise mutual information:**

- How much more do events x and y co-occur than if they were independent?

$$\text{PMI}(X, Y) = \log_2 \frac{P(x, y)}{P(x)P(y)}$$

55



Pointwise Mutual Information

- **Pointwise mutual information:**

- How much more do events x and y co-occur than if they were independent?

$$\text{PMI}(X, Y) = \log_2 \frac{P(x, y)}{P(x)P(y)}$$

- **PMI between two words:**

- How much more do two words co-occur than if they were independent?

$$\text{PMI}(\text{word}_1, \text{word}_2) = \log_2 \frac{P(\text{word}_1, \text{word}_2)}{P(\text{word}_1)P(\text{word}_2)}$$

56



How to Estimate Pointwise Mutual Information

- Query search engine (Altavista)

- $P(\text{word})$ estimated by $\text{hits}(\text{word}) / N$
- $P(\text{word}_1, \text{word}_2)$ by $\text{hits}(\text{word}_1 \text{ NEAR } \text{word}_2) / N^2$

$$\text{PMI}(\text{word}_1, \text{word}_2) = \log_2 \frac{\text{hits}(\text{word}_1 \text{ NEAR } \text{word}_2)}{\text{hits}(\text{word}_1)\text{hits}(\text{word}_2)}$$

57



Does phrase appear more with “poor” or “excellent”?

$$\text{Polarity}(\text{phrase}) = \text{PMI}(\text{phrase}, \text{"excellent"}) - \text{PMI}(\text{phrase}, \text{"poor"})$$

$$= \log_2 \frac{\text{hits}(\text{phrase NEAR "excellent"})}{\text{hits}(\text{phrase})\text{hits}(\text{"excellent"})} - \log_2 \frac{\text{hits}(\text{phrase NEAR "poor"})}{\text{hits}(\text{phrase})\text{hits}(\text{"poor"})}$$

$$= \log_2 \frac{\text{hits}(\text{phrase NEAR "excellent"})}{\text{hits}(\text{phrase})\text{hits}(\text{"excellent"})} - \frac{\text{hits}(\text{phrase})\text{hits}(\text{"poor"})}{\text{hits}(\text{phrase NEAR "poor"})}$$

$$= \log_2 \frac{\text{hits}(\text{phrase NEAR "excellent"})\text{hits}(\text{"poor"})}{\text{hits}(\text{phrase NEAR "poor"})\text{hits}(\text{"excellent"})}$$

58



Phrases from a thumbs-up review

Phrase	POS tags	Polarity
online service	JJ NN	2.8
online experience	JJ NN	2.3
direct deposit	JJ NN	1.3
local branch	JJ NN	0.42
...		
low fees	JJ NNS	0.33
true service	JJ NN	-0.73
other bank	JJ NN	-0.85
inconveniently located	JJ NN	-1.5
Average		0.32

59



Phrases from a thumbs-down review

Phrase	POS tags	Polarity
direct deposits	JJ NNS	5.8
online web	JJ NN	1.9
very handy	RB JJ	1.4
...		
virtual monopoly	JJ NN	-2.0
lesser evil	RBR JJ	-2.3
other problems	JJ NNS	-2.8
low funds	JJ NNS	-6.8
unethical practices	JJ NNS	-8.5
Average		-1.2

60



Results of Turney algorithm

- 410 reviews from Epinions
 - 170 (41%) negative
 - 240 (59%) positive
- Majority class baseline: 59%
- Turney algorithm: 74%
- Phrases rather than words
- Learns domain-specific information

61



Summary on Learning Lexicons

- Advantages:
 - Can be domain-specific
 - Can be more robust (more words)
- Intuition
 - Start with a seed set of words ('good', 'poor')
 - Find other words that have similar polarity:
 - Using "and" and "but"
 - Using words that occur nearby in the same document
 - Using WordNet synonyms and antonyms (be careful with senses)

62



Sentiment Analysis

Other Sentiment Tasks

63



Finding sentiment of a sentence

- Important for finding aspects or attributes
 - Target of sentiment
- The food was great but the service was awful

64



Finding aspect/attribute/target of sentiment

M. Hu and B. Liu. 2004. Mining and summarizing customer reviews. In Proceedings of KDD.
S. Blair-Goldensohn, K. Hannan, R. McDonald, T. Neylon, G. Reis, and J. Reynar. 2008. Building a Sentiment Summarizer for Local Service Reviews. WWW Workshop.

- Frequent phrases + rules
 - Find all highly frequent phrases across reviews ("fish tacos")
 - Filter by rules like "occurs right after sentiment word"
 - "...great fish tacos" means fish tacos a likely aspect

Casino	casino, buffet, pool, resort, beds
Children's Barber	haircut, job, experience, kids
Greek Restaurant	food, wine, service, appetizer, lamb
Department Store	selection, department, sales, shop, clothing

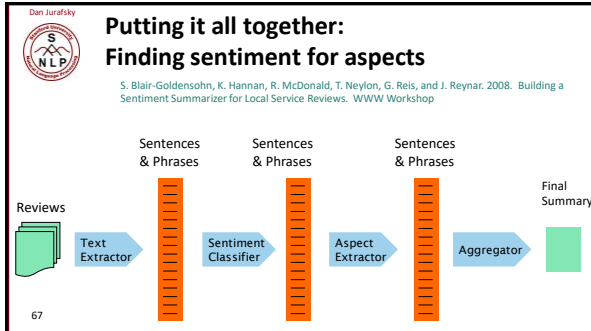
65



Finding aspect/attribute/target of sentiment

- The aspect name may not be in the sentence
- For restaurants/hotels, aspects are well-understood
- Supervised classification
 - Hand-label a small corpus of restaurant review sentences with aspect
 - food, décor, service, value, NONE
 - Train a classifier to assign an aspect to a sentence
 - "Given this sentence, is the aspect *food*, *décor*, *service*, *value*, or *NONE*?"

66



Results of Blair-Goldensohn et al. method

Rooms (3/5 stars, 41 comments)

- (+) The room was clean and everything worked fine – even the water pressure ...
- (+) We went because of the free room and was pleasantly pleased ...
- (-) ...the worst hotel I had ever stayed at ...

Service (3/5 stars, 31 comments)

- (+) Upon checking out another couple was checking early due to a problem ...
- (+) Every single hotel staff member treated us great and answered every ...
- (-) The food is cold and the service gives new meaning to SLOW.

Dining (3/5 stars, 18 comments)

- (+) our favorite place to stay in biloxi.the food is great also the service ...
- (+) Offer of free buffet for joining the Play

68

Methods work best when classes have equal frequencies

- If not balanced (common in the real world)
 - can't use accuracies as an evaluation
 - need to use F-scores
- Severe imbalancing also can degrade classifier performance
- Two common solutions:
 - Resampling in training
 - Random undersampling
 - Cost-sensitive learning
 - Penalize SVM more for misclassification of the rare thing

69

How to deal with 7 stars?

Bo Pang and Lillian Lee. 2005. Seeing stars: Exploiting class relationships for sentiment categorization with respect to rating scales. ACL, 115–124

- Map to binary
- Use linear or ordinal regression
 - Or specialized models like metric labeling

70

Summary on Sentiment

- Generally modeled as classification / regression task
 - predict a binary or ordinal label
- Features:
 - Negation is important
 - Using all words (in naive bayes) works well for some tasks
 - Finding subsets of words may help in other tasks
 - Hand-built or automated polarity lexicons
- Aspect extraction

71

Recall: Scherer Typology of Affective States

- Emotion:** brief organically synchronized ... evaluation of a major event
 - angry, sad, joyful, fearful, ashamed, proud, elated
- Mood:** diffuse non-caused low-intensity long-duration change in subjective feeling
 - cheerful, gloomy, irritable, listless, depressed, buoyant
- Interpersonal stances:** affective stance toward another person in a specific interaction
 - friendly, flirtatious, distant, cold, warm, supportive, contemptuous
- Attitudes:** enduring, affectively colored beliefs, dispositions towards objects or persons
 - liking, loving, hating, valuing, desiring
- Personality traits:** stable personality dispositions and typical behavior tendencies
 - nervous, anxious, reckless, morose, hostile, jealous

72



Computational work on other affective states

- **Emotion:**
 - Detecting annoyed callers to dialogue system
 - Detecting confused/frustrated versus confident students
- **Mood:**
 - Finding traumatized or depressed writers
- **Interpersonal stances:**
 - Detection of flirtation or friendliness in conversations
- **Personality traits:**
 - Detection of extroverts

73



Detection of Friendliness

Ranganath, Jurafsky, McFarland

- Friendly speakers use collaborative conversational style
 - Laughter
 - Less use of negative emotional words
 - More sympathy
 - That's too bad I'm sorry to hear that
 - More agreement
 - I think so too
 - Less hedges
 - kind of sort of a little ...

74