

# MIE 1513 Decision Support Systems

## Lab and Assignment 4:

### Data Science and NLP for Customer Review Analysis

March 11, 2019

This lab and assignment involves performing analysis of real hotel review data crawled from the Tripadvisor website to automatically identify positive and negative keywords and phrases associated with hotels and to better understand characteristics of data analysis tools, extracting explanatory review summaries, and human reviewing behavior.

- Programming language: Python (Jupyter IPython Environment)
- Due Date: Posted in Syllabus

**Marking scheme and requirements:** Full marks will be given for (1) working, readable, reasonably efficient, documented code that achieves the assignment goals and (2) for providing appropriate answers to the questions in a Jupyter notebook `ds-assignment.ipynb` committed to the student's assignment repository.

Please adhere to the collaboration policy on the course website – people you discussed the assignment solution with, or websites with source code you used should be listed in the submitted Jupyter notebook.

#### What/how to submit your work:

- All your code should be included in a notebook named `ds-assignment.ipynb`.
- Commit and push your work to your github repository in order to submit it. Your last commit and push before the assignment deadline will be considered to be your submission. You can check your repository online to make sure that all required files have actually been committed and pushed to your repository.
- A link to create a personal repository for this assignment is posted on QUERCUS.

# 1 Before and in the Introductory lab

In this lab, we will familiarize ourselves with the **nlTK** library and an nltk sentiment analysis package called **Vader**. Please familiarize yourself with this tool through the examples you will be taught by in the lab.

In lab, we will provide some examples of the types of data analysis required for the assignment, however you will be required to extend these analyses to additional cases.

## 2 Tripadvisor Crawler

In this assignment, all students will be required to crawl a specific dataset for hotel reviews in a city that should not overlap with any of your classmates.

To use the Tripadvisor crawler, please download the Python files from <https://github.com/aesuli/trip-advisor-crawler>. First run the crawler `trip-advisor-crawler.py` from the command with your own city selection and then parse all crawled html file into csv using `trip-advisor-parser.py` before loading into your DSVM environment. More details provided below.

### 2.1 Command-line Examples

If you search the hotel reviews of city Niagara Falls on Tripadvisor, you will see the web address is [https://www.tripadvisor.ca/Hotels-g154998-Niagara\\_Falls\\_Ontario-Hotels.html](https://www.tripadvisor.ca/Hotels-g154998-Niagara_Falls_Ontario-Hotels.html). From the web address, you will find two important pieces of information for the command line invocation. First, the domain is 'ca'. And second, the city code of the Niagara Falls is **154998** on Tripadvisor.

To download the reviews, you will also need to specify the local path for the files. In this example, the location is the "data" folder under your assignment folder. Then you can extract the reviews using the following command line:

```
$python3 trip-advisor-crawler.py -o data ca:154998
```

The downloaded reviews are in **html** format. You still need to process the data into the **csv** format so that we can easily load it into Python. Fortunately, you do NOT need to write the code yourself. Please run the following command line after extraction from previous step:

```
$python3 trip-advisor-parser.py -d data -o reviews.csv
```

**Warning:** You may not be able to crawl all the hotels within a city; stop when you have crawled for more than 10 hours (i.e., run your crawler overnight). Be aware that there are two additional things you need to do if you stop crawling manually or the file cannot be parsed due to an incomplete crawl:

- In `ids.txt` file, delete the last hotel's reviews entirely (all reviews for that hotel).
- In the data folders, delete the folder of the same hotel entirely.

You may crawl a city that has less than 20 hotels listed on Tripadvisor, but please do not crawl anything significantly smaller; this may lead to insufficient data to get reasonable answers for the assignment.

### 3 Main Assignment

Please answer the questions below in an IPython notebook that you must submit via github. In the following, ground truth rating (star rating) can be binarized; if so, explain your rationale for how your binarize.

#### Q1. Sentiment Analysis and Aggregation

- (a) Compute average Vader sentiment and average ground truth rating per hotel.
- (b) Rank hotels by
  - (i) Average Ground Truth Sentiment
  - (ii) Average Vader Compound Sentiment Score

Show both top-5 and bottom-5 for both ranking methods. Do they agree or are there interesting differences?

#### Q2. Frequency Analysis

- (a) Use term frequency of the words for (i) positive reviews and (ii) negative with ground truth sentiment to rank the top-50 most frequent non-stopwords in the review collection. Do you note anything interesting and/or locale-specific about these top-ranked words?
- (b) Repeat this analysis for the top-50 noun phrases and note any interesting results.

#### Q3. Mutual Information

- (a) Use mutual information (MI) with ground truth sentiment to rank the top-50 most sentiment-bearing non-stopwords in the review collection. Do you note anything interesting and/or locale-specific about these top-ranked words?
- (b) Repeat this analysis for the top-50 noun phrases and note any interesting results.

#### Q4. Pointwise Mutual Information

- (a) For ground truth sentiment, calculate the top-50 words according to PMI of the word occurring with (i) positive reviews and (ii) negative reviews. Do you note anything interesting and/or locale-specific about these top-ranked words?
- (b) Repeat this analysis for the top-50 noun phrases and note any interesting results.
- (c) Repeat this analysis for the single top **and** single bottom hotel (according to the ground truth rating). Do you gain any useful hotel-specific insights about what is good and bad about these two hotels? If not, explain why not.

#### Q5. General Plots

- (a) Histogram
  - (a) Show separate histograms of ground truth and Vader sentiment scores (ignore hotel ID). Do you notice any interesting differences?

- (b) Show a histogram of the number of reviews per hotel. Do you notice any interesting trends? Are these expected?
- (b) Boxplots
- (a) In two plots, one for ground truth star rating and one for Vader sentiment, show a plot of 5 side-by-side boxplots of these scores.
  - (b) Report the mean and variance of the ground truth and Vader sentiment scores for the top-5 ranked hotels according to star rating.
  - (c) Which do you find more informative, the boxplots or the mean and variance, or are they equally informative? Why?
- (c) Scatterplots and heatmaps
- (a) Show both a scatterplot **and** heatmap of ground truth score (star rating) versus Vader sentiment score. Each review is a point on the scatterplot. Do you notice anything interesting? What does this tell you about star ratings vs. Vader sentiment scores? What does this tell you about human ratings and/or Vader sentiment analysis?
  - (b) Show two scatterplots **and** two heatmaps of the length of reviews versus each of ground truth score and Vader sentiment score. Each review is a point on the scatterplot. Are there any trends?
  - (c) Show two scatterplots of the number of reviews per hotel versus each of average ground truth score and average Vader sentiment score. In this case, each hotel is a single point on the scatterplot. Are there any trends?

## 4 Clarifications

- This assignment is graded entirely via code review – your **submitted** IPython notebook must contain all experimental output – notebooks with empty output cells will **not** be graded.
- Please do not commit/push the crawled data to your github Repository (github is not for data storage) – it should only contain your IPython notebook.
- Some crawled reviews do not contain any words/strings. You should perform data preprocessing to remove those rows.
- Before you use a Pandas Dataframe for merge and join operations, please first read and understand the respective function descriptions in the online official Pandas documentation.